

Realistic Time-Domain Synthesis of Gravitational-Wave Detector Glitches using Class-Conditional Derivative Generative Adversarial Networks

Tom Dooney^{1,2}, Mees de Boer², Harsh Narola^{1,2}, Melissa Lopez^{1,2},
Stefano Bromuri³, Daniel Stanley Tan³, and Chris Van Den Broeck^{1,2}

¹*Nikhef, Science Park 105, 1098 XG Amsterdam, The Netherlands*

²*Institute for Gravitational and Subatomic Physics (GRASP),*

Utrecht University, Princetonplein 1, 3584 CC Utrecht, The Netherlands and

³*Faculty of Science, Open Universiteit, Valkenburgerweg 177, 6419 AT Heerlen, The Netherlands*

(Dated: June 2, 2026)

Gravitational-wave detectors such as LIGO, Virgo, and KAGRA are highly sensitive instruments susceptible to numerous noise sources. Short-duration transient noise events, known as glitches, pose a particular challenge for data analysis pipelines as they can mimic or obscure astrophysical signals. We present GlitchGAN, a class-conditional generative model, built upon the Conditional Derivative GAN (cDVGAN) architecture, that is capable of synthesizing realistic glitches directly in the time domain. The model is trained on high-quality reconstructions of seven common glitch types observed during LIGO’s third observing run (O3): Blip, Fast Scattering, Koi Fish, Low-Frequency Burst, Scattered Light, Tomte, and Whistle. We show that GlitchGAN generalizes effectively, learning to reproduce a diverse and physically consistent glitch space directly from these reconstructions. Moreover, because the model is conditioned on glitch class, it can generate *hybrid* or transitional glitch morphologies by interpolating across the class-conditioning vector after training. Synthetic glitches are validated against real glitches using the Gravity Spy classifier, widely used in the GW community for glitch classification, and an unsupervised analysis using UMAP embeddings. Gravity Spy classifies the majority of GlitchGAN’s synthetic glitches as the correct class while the UMAP analysis shows substantial overlap between real and synthetic samples in the reduced latent space. We further highlight a critical limitation of magnitude-only spectrograms: classifiers operating on magnitude Q -transforms can confidently misclassify physically unrealistic glitches from less robust models, underscoring the need for complementary validation methods that preserve phase information.

I. INTRODUCTION

The direct detection of gravitational waves (GWs) by the LIGO [1] and Virgo [2] interferometers has opened a new observational window on the universe, enabling measurements of compact binary coalescences and other cataclysmic astrophysical events [3]. The extreme sensitivity required for these detections, however, renders the instruments equally susceptible to instrumental and environmental noise sources [4]. Among these, short-duration transient noise artefacts known as *glitches* [5–7] represent a persistent challenge for data analysis pipelines. Glitches exhibit diverse morphologies—ranging from impulsive millisecond-duration events to extended, dispersive features spanning several seconds—and can mimic or obscure astrophysical signals [8–10]. When a glitch overlaps with a genuine GW signal it can significantly bias parameter estimation [11–13], and approximately 25 of the ~ 90 confident detections up to the end of O3 required some form of glitch mitigation [3, 11, 14–17]. With the anticipated sensitivity improvements of next-generation detectors such as Einstein Telescope [18] and Cosmic Explorer [19], the rate and impact of glitches are expected to increase substantially [20].

The ability to synthesize accurate and realistic glitch populations is therefore of growing importance. Such capability enhances the realism of detector simulations and supports software injection campaigns [6, 21–23], enabling more rigorous validation of GW detection

pipelines and parameter estimation algorithms. Realistic synthetic glitches are also essential for the construction of mock data challenges (MDCs) and for data augmentation in machine learning-based glitch classification systems. A possible approach is to model glitches mathematically, given *a priori* knowledge of their production mechanism [24]. Other works have developed models that mimic their principal characteristics using a few parameters [25]. However, such approaches either require prior knowledge of the glitch’s physical origin, which isn’t always available, or are limited to a restricted parametric form, making them difficult to scale across the full diversity of observed morphologies.

Recently, generative models have been applied broadly within GW data analysis to address multiple challenges in data synthesis [26, 27]. For glitch synthesis, Generative Adversarial Networks (GANs) [28] have been used to augment spectrogram-based datasets for downstream classification tasks [29, 30]. The *gengli* glitch generator [31, 32] demonstrated that a GAN trained on Blip glitches extracted from detector data using BayesWave [33], the state-of-the-art method for reconstructing glitches from detector data, can successfully learn their underlying time-domain distribution, providing a useful framework for simulations and MDCs. McGANn [34] implemented a conditional GAN capable of synthesizing five distinct waveform classes analogous to GW bursts in one model, with class interpolation enabling the generation of hybrid morphologies that blend features from the learned

classes. Powell et al. [30] implemented generative modelling on a broader glitch class space by training separate GANs on magnitude Q -scan representations of each Gravity Spy class. However, extending this holistic approach to a single model operating directly in the time domain remains unexplored, despite its advantages in preserving phase information and enabling direct use of time-series data in MDCs and detection pipelines.

The principal obstacle to realistic time-domain glitch synthesis has been the availability of high-quality training data. BayesWave [33] is computationally expensive, and converges unreliably for certain glitch morphologies, making it impractical to construct the large-scale, diverse datasets that cover the broad glitch space. The recent development of DeepExtractor [35]—a deep learning U-Net framework for time-domain glitch reconstruction—offers fast, high-fidelity reconstructions that preserve both amplitude and phase information without imposing waveform smoothness assumptions, achieving a $10,000\times$ speed-up over BayesWave in realistic simulations.

In this paper, we introduce GlitchGAN, a class-conditional model trained on a dataset of DeepExtractor reconstructions of seven common glitch classes observed during LIGO’s O3 run: Blip, Fast Scattering, Koi Fish, Low-Frequency Burst, Scattered Light, Tomte, and Whistle. GlitchGAN inherits the cDVGAN architecture [36, 37], which features an auxiliary discriminator applied to the time derivative of signals, providing additional regularisation that improves training stability and sample fidelity compared to standard GAN counterparts.

We validate the quality of the DeepExtractor training data using the Gravity Spy classifier [38, 39] and unsupervised UMAP embeddings [40], establishing a baseline against which synthetic glitches can be assessed. Applying the same evaluation to GlitchGAN’s outputs, we show that the model captures much of the morphological diversity across all seven classes, with substantial overlap between real and generated distributions in the UMAP latent space and with the majority of synthetic glitches correctly identified by Gravity Spy. We also demonstrate that GlitchGAN can interpolate between classes, producing novel hybrid morphologies that mix features from the learned classes. We further highlight a fundamental limitation of magnitude-only time-frequency representations as validation tools: when applied to physically unrealistic glitches from a diffusion-based comparison model [41], Gravity Spy assigns correct class labels with high confidence despite clear inconsistency in the time domain, because the Q -transform discards phase information entirely. This underscores the necessity of complementary, phase-preserving validation methods.

The remainder of the paper is organized as follows: Section II describes the curation and preparation of our training dataset. Section III details the GlitchGAN architecture and training procedure. Section IV presents validation results across multiple metrics. Finally, Section V discusses the implications of our findings and opportunities for future work.

II. DATA

A. Data curation

We construct our training dataset using DeepExtractor reconstructions of real LIGO glitches. To ensure high-quality, morphologically representative samples, we apply strict selection criteria:

1. Gravity Spy classification confidence $\sigma_{\text{GS}} \geq 0.9$, ensuring that the reconstructed glitches are representative of well-characterized classes.
2. Signal-to-noise ratio (SNR) ≥ 15 (sufficient amplitude to minimize reconstruction artifacts)
3. Detectors: LIGO Hanford (H1) and Livingston (L1)
4. Observing runs: O3a and O3b

We include seven glitch classes in the training of GlitchGAN, namely; Blip, Fast Scattering, Koi Fish, Low-Frequency Burst, Scattered Light, Tomte, and Whistle.

Blip glitches have a characteristic morphology of a symmetric teardrop shape in time-frequency in the range [30, 300] Hz with short-durations, ~ 0.04 s. They appear in both Hanford and Livingston, as well as Virgo and GEO 600 [42]. Due to their abundance and form, these glitches hinder both the unmodeled burst and modelled CBC searches, with particular emphasis on compact binaries with large total mass, highly asymmetric component masses, and spins anti-aligned with the orbital angular momentum [10].

Fast Scattering glitches appear as short-duration arches ($\sim 0.2 - 0.3$ s) in the frequency range [20, 60] Hz. These glitches are strongly correlated with ground motion in range [0.1 – 0.3] Hz and [1 – 6] Hz, which in turn is associated with thunderstorms and human activity near the detector. Note that this class was added during O3, and it is more abundant in Livingston than Hanford, due to differences in ground motion and detector sensitivity [43].

Koi Fish glitches get their name from their frontal resemblance to Koi Fishes. They are also similar to *Blips* but typically feature a higher SNR, spanning the frequency range of [20, 1000] Hz.

Low-Frequency Bursts are short-duration glitches (~ 0.25 s) in the frequency range [10 – 20] Hz with a distinctive blob shape. These occurrences were prevalent in Livingston data during O1 and Hanford data in O3a.

Scattered Light glitches have longer duration harmonics ($\sim 2.0 - 4.0$ s), and in the time-frequency domain, they appear as arches being often stacked on top of each other. These glitches are problematic since their frequency content lies in the band of interest of GW astrophysical events. In O3, they were found to be coupled with the relative motion between the optical suspension system’s end test-mass chain and the reaction-mass chain [43].




FIG. 1: Comparison of real and GlitchGAN-generated glitches across the seven glitch classes used for training, shown in both representations. **Row 1 (Real, time domain):** DeepExtractor reconstructions of real glitches. **Row 2 (Real, Q-transform):** Corresponding Q-scans of the same real glitches. **Row 3 (Synthetic, time domain):** Glitches generated by GlitchGAN. **Row 4 (Synthetic, Q-transform):** Corresponding Q-scans of the generated glitches. Each column shows one representative sample per class.

Tomte glitches are also short-duration (~ 0.25 s) with characteristic triangular morphology.

Whistles have a characteristic V, U or W shape at higher frequencies with typical durations ~ 0.25 s. They are caused when radio-frequency signals beat with the voltage-controlled oscillators [44].

These seven classes collectively span wide time and frequency ranges, encompassing much of the morphological variability relevant to transient instrumental noise. While it is a common glitch class, we exclude the Extremely Loud class because its characterization is dominated by exceptionally high SNR rather than distinctive morphology, making it a poor target for generative modeling.

B. Dataset Preprocessing

We reconstruct with DeepExtractor a 2-second window centered on the glitch trigger time at a sampling rate of 4096 Hz, yielding 8192 samples per glitch waveform. Our goal was to obtain approximately 5,000 samples per class (1,250 per detector per observing run) under the selection criteria above. Unfortunately, data availability constraints limited acquisition for some classes. Fast Scattering yielded only 2,774 samples, and Low-Frequency Burst provided 4,316 samples. The total curated dataset contained 31,511 samples.

To mitigate training instabilities from class imbalance, we applied bootstrap resampling (sampling with replace-

ment) to oversample all classes to 5,000 samples each, yielding a balanced training set of 35,000 total samples. Dataset composition by detector, observing run, and class is detailed in Table I. Before training, all reconstructed time series are normalized to the range $[-1, 1]$, followed by subtracting the mean. This normalization constrains the generative model to operate at a single effective SNR scale, but generated samples can be rescaled post-generation to any desired SNR for downstream applications.

TABLE I: Composition of the curated glitch dataset. Number of samples per class, detector (H1: Hanford, L1: Livingston), and observing run (O3a, O3b). Classes are bootstrapped to 5000 samples each for balanced training.

Glitch Class	H1		L1	
	O3a	O3b	O3a	O3b
Blip	1246	1247	1245	1243
Fast Scattering	143	143	1244	1244
Koi Fish	1247	1248	1243	1247
Low Freq. Burst	1246	1246	912	912
Scattered Light	1236	1243	1246	1248
Tomte	1031	1031	1247	1246
Whistle	1232	1232	1231	1232

III. METHODS

A. Model architecture: cDVGAN

Generative adversarial networks (GANs) [28], wherein a generator and discriminator are trained in an adversarial min-max game, have established themselves as a powerful framework for learning complex high-dimensional data distributions. Wasserstein GANs (WGANs) [45] improve training stability by replacing the Jensen–Shannon divergence with the Wasserstein-1 distance, and the gradient penalty variant (WGAN-GP) [46] further stabilizes optimization by enforcing a Lipschitz constraint on the discriminator. Conditional extensions [47] allow the generator to be steered towards a target class, enabling multi-class synthesis within a single model.

cDVGAN [36] is a conditional WGAN-GP that augments the standard two-player adversarial game with an auxiliary discriminator applied to the first-order time derivative of signals. Here we summarise the key components; full layer-by-layer architecture details are shown in Appendix A.

1. Training objective

Under the WGAN-GP framework, the loss for discriminator D is

$$L_D = \mathbb{E}_{\tilde{x}}[D(\tilde{x}, c)] - \mathbb{E}_x[D(x, c)] + \lambda \mathbb{E}_{\tilde{x}} \left[(\|\nabla_{\tilde{x}} D(\tilde{x}, c)\|_2 - 1)^2 \right], \quad (1)$$

where x and $\hat{x}$ are real and generated samples respectively, $\tilde{x}$ is a random interpolation between them, and $\lambda = 10$ is the gradient-penalty weight. The generator is updated to minimise

$$L_G = -\mathbb{E}_{\hat{x}}[D(\hat{x}, c)]. \quad (2)$$

In cDVGAN a second discriminator D_∂ operates on the first-order derivative $\dot{x} = dx/dt$. The combined generator loss is

$$L_G^{\text{cDVGAN}} = -\frac{1}{2} \left(\mathbb{E}_{\hat{x}}[D(\hat{x}, c)] + \mathbb{E}_{\dot{\hat{x}}}[\dot{D}(\dot{\hat{x}}, c)] \right), \quad (3)$$

where both discriminators contribute equally. Each discriminator is updated five times per generator update.

2. Generator

The generator G maps a noise vector $z \sim \mathcal{N}(0, I_{100})$ and a one-hot class vector $c \in \mathbb{R}^7$ to a time series of length 8192. The class vector is first projected to a 32-dimensional embedding and concatenated with z , then passed through a dense layer and reshaped before five successive upsample-and-convolve blocks with batch normalisation and ReLU activations expand the sequence to the target length.

3. Discriminators

Both D and D_∂ are 1D convolutional networks that map an input sequence and a class vector to a scalar critic value. Class conditioning uses projection [47]: a learned class embedding is dot-producted with the discriminator’s feature vector and added to the scalar output. D operates directly on the 8192-sample waveform; D_∂ operates on the 8191-sample first-order difference sequence $\dot{x} = x_{t+1} - x_t$. The derivative discriminator provides additional regularisation by penalising unphysical high-frequency artefacts in generated waveforms [37].

B. Training procedure

The model was trained for 500 epochs with a batch size of 64 using the RMSprop optimiser and a learning rate of 10^{-4} for all network components. Training was performed on a NVIDIA A100 GPU (approximately 12 hours). Although training was conducted for the full 500 epochs, the final GlitchGAN configuration uses the model weights saved at epoch 210.

The selection of epoch 210 was based on a combination of qualitative and quantitative assessments. Generated glitches were inspected throughout training, and epoch 210 was found to produce the most realistic glitch morphologies. This choice is further supported by the generator loss evolution (Appendix B), which exhibits relatively stable behaviour around this epoch compared to other stages of training. Additionally, the UMAP analysis presented in Section IV B 2 shows improved overlap between the distributions of real and generated glitches at epoch 210 relative to models saved at epochs 100, 200, 300, 400, and 500.

No hyperparameter tuning was performed for this dataset; the architecture and hyperparameters are identical to those validated in [36], allowing us to assess out-of-the-box generalisation to this more complex and realistic glitch population. The implementation uses TensorFlow [48] and is publicly available at <https://github.com/tomdooney95/glitchgan>.

C. Comparison with a diffusion-based model

As a comparison against our cDVGAN architecture in GlitchGAN, we also trained another type of generative model on our dataset known as a latent Diffusion Transformer (DiT), which was adapted for 1D time series. Diffusion models [41, 49, 50] learn to generate data by explicitly modelling the process of iteratively adding and then removing noise from samples drawn from a simple prior distribution. They have become increasingly popular in recent years for their ability to generate realistic synthetic data. Full-resolution diffusion proved computationally prohibitive, so we first trained a convolutional autoencoder to compress inputs from (8192, 1) to

a latent space of (256, 16), reducing computational cost by approximately three orders of magnitude while preserving essential morphological structure. The DiT was then trained in this latent space, with generated samples decoded back to the time domain for evaluation. This model serves primarily to illustrate the limitation of magnitude-only validation metrics, as discussed in Section IV.

D. Evaluation Methodology

We evaluate both the training data (DeepExtractor reconstructions) and GlitchGAN-generated glitches using two complementary approaches: classification with the Gravity Spy classifier and visualisation of low-dimensional embeddings via t-SNE and UMAP. For DeepExtractor, this establishes a quantitative baseline confirming the training data is morphologically representative; for GlitchGAN, it assesses how faithfully the model reproduces the underlying glitch distributions.

1. Gravity Spy Classification

We reconstruct with DeepExtractor 100 randomly selected glitches from the Gravity Spy dataset according to the criteria in Section II, drawing 25 samples per detector (H1, L1) per observing run (O3a, O3b). Since Gravity Spy classifies glitches surrounded by detector background noise i.e. $g(t) + n(t)$, each DeepExtractor reconstruction $\hat{g}(t)$ is injected into a segment of ‘quiet’ Hanford detector noise from O3 prior to reclassification (representing $n(t)$).

We follow a similar approach to evaluate the GlitchGAN-generated synthetic glitches, injecting them into the same noise segments before classification. The only difference is that we have to explicitly rescale the generated glitches to a fixed SNR of 50 to ensure they are sufficiently loud for classification, since the generator is trained on normalized waveforms and does not learn an explicit SNR distribution.

2. Dimensionality Reduction

To visualise and compare the high-dimensional feature space of real and synthetic glitches, we use two complementary nonlinear dimensionality reduction methods: *t-distributed Stochastic Neighbor Embedding* (t-SNE) [51] and *Uniform Manifold Approximation and Projection* (UMAP) [52]. While *Principal Component Analysis* (PCA) captures global variance through linear projections, it cannot resolve the nonlinear structure present in morphologically diverse glitch populations; t-SNE and UMAP are instead designed to preserve local cluster structure in the embedding. t-SNE embeds data by minimising the Kullback–Leibler divergence between proba-

bility distributions encoding pairwise similarities in the original and reduced spaces; it excels at revealing local cluster structure but can distort global geometry. UMAP constructs a nearest-neighbour graph in the high-dimensional space and optimises a low-dimensional counterpart to preserve both local and global topological structure, scaling efficiently to the dataset sizes used here.

For DeepExtractor’s reconstructions, we aim to verify that the reconstructed glitches cluster according to their correct Gravity Spy class in an unsupervised manner, confirming that the reconstructions preserve the morphological characteristics recognized by Gravity Spy. For GlitchGAN-generated synthetic glitches, we evaluate whether they cluster with the real reconstructions of their respective classes, indicating that the model has successfully learned to reproduce the underlying morphological distribution of each class.

IV. RESULTS

A. Validation of DeepExtractor training data

1. Classification evaluation using Gravity Spy

Tables II and III summarize the results of classifying 700 DeepExtractor reconstructions (100 per class) with Gravity Spy. Since Gravity Spy supports additional classes beyond our seven training classes, the matrix is asymmetric. Of the 700 reconstructed glitches, 627 are correctly identified by Gravity Spy, corresponding to 89.57% overall accuracy. Correctly classified reconstructions achieve a mean confidence of 96.77%, confirming that DeepExtractor effectively preserves the morphological characteristics recognized by Gravity Spy. Misclassified samples show lower confidence of 84.84%, suggesting that classifier confidence may serve as a useful quality veto.

TABLE II: Summary of Gravity Spy results on DeepExtractor reconstructions (100 samples per class, 700 total).

Total samples	700
Correct predictions	627
Incorrect predictions	73
Overall accuracy	89.57 %
Mean confidence (correct)	96.77 %
Mean confidence (incorrect)	84.84 %

Table III shows that class-level accuracy varies considerably. All Blip and Low-Frequency Burst glitches are correctly classified, while Fast Scattering and Koi Fish achieve the lowest accuracies at 77% and 79%, respectively. Figure 2 shows the full confusion matrix. Most misclassifications occur between morphologically similar classes: among the 21 misclassified Koi Fish glitches, 14




FIG. 2: Confusion matrix for Gravity Spy classification of 100 DeepExtractor reconstructions per class (700 total). Each cell shows the number of predictions with the mean Gravity Spy confidence in brackets.

TABLE III: Per-class correct prediction count and mean Gravity Spy confidence for DeepExtractor reconstructions.

True label	Correct count	Mean confidence (%)
Blip	100	99.30
Fast Scattering	77	89.78
Koi Fish	79	91.59
Low Frequency Burst	100	99.85
Scattered Light	83	96.95
Tomte	96	98.21
Whistle	92	99.33

were labelled as Extremely Loud, consistent with their exceptionally high SNRs (220–1270) overwhelming finer morphological details. The 23 misclassified Fast Scattering samples are distributed across several related classes (Scattered Light, Tomte, Repeating Blips), reflecting the known morphological similarity between low-frequency scattering noise types. Overall, the misclassification patterns are morphologically explainable and do not indicate systematic failure of the reconstruction.

2. Unsupervised clustering with t-SNE and UMAP

Figure 3 shows three-dimensional t-SNE (left) and UMAP (right) embeddings of the DeepExtractor reconstructions, visualized from three complementary viewpoints. The reconstructions form well-separated clusters that align closely with the Gravity Spy labels, demonstrating that morphologically distinct glitch types occupy distinct regions of the time-domain latent space. Partial overlap between the Blip and Koi Fish clusters is visible, consistent with prior observations that these classes may share a physical origin [53]. The good clustering structure validates that the 31,511 DeepExtractor recon-

structions used to train GlitchGAN are morphologically representative of seven distinct classes and are not dominated by reconstruction artifacts or label noise.

B. Validation of GlitchGAN synthetic glitches

1. Classification evaluation using Gravity Spy

To quantify synthetic glitch realism we follow a similar approach to the DeepExtractor evaluation, injecting 100 generated glitches per class into quiet Hanford O3 background noise and classifying them with Gravity Spy. The resulting confusion matrix shown in Figure 4 reflects classifier predictions and confidence scores in brackets. The overall classification accuracy is $\sim 72\%$. While this may appear low, the misclassifications are consistent with known limitations of Gravity Spy and the use of a fixed injected SNR of 50.

Most misclassifications occur in the Koi Fish and Scattered Light classes. While only 20% of the generated Koi Fish samples were correctly classified, the majority of the remaining samples were misidentified as Blip, Blip Low Frequency, or Tomte. This behaviour is consistent with known morphological similarities between Koi Fish and these glitch classes, which may arise from related physical mechanisms differing primarily in SNR [53]. Moreover, increasing the injected SNR of the generated Koi Fish samples to 150 raises the classification accuracy to 79% with an average confidence of $\sim 95\%$, further illustrating the dependence of Gravity Spy classifications on the injected SNR.

This behaviour is further supported by injecting glitches at an SNR of 20. In this lower-SNR regime, none of the generated Koi Fish samples were correctly classified. Instead, they were misidentified as other morphologically similar classes. A comparable SNR dependence




FIG. 3: Three-dimensional *t*-SNE (left) and UMAP (right) embeddings of DeepExtractor reconstructions across three viewpoints. Colours correspond to Gravity Spy glitch classes. Well-separated clusters demonstrate that reconstructions preserve class-specific morphological structure in the time domain. The partial overlap between Blip and Koi Fish (visible in the middle row) is consistent with their known morphological proximity.

is observed for the Whistle class. Injecting the generated Whistle glitches at an SNR of 50 yields an accuracy of 94%, whereas injecting them at an SNR of 20 leads to frequent misclassification as No Glitch.

For Scattered Light, Gravity Spy often predicts Fast Scattering, which is expected given their related scattering origins at low frequency. This was also observed for DeepExtractor’s original reconstructions of both Fast

Scattering and Scattered Light shown in Figure 2.

Thus, if Blip, Tomte, and Blip Low Frequency are regarded as acceptable misclassifications for the Koi Fish class, and Fast Scattering is considered an acceptable misclassification for Scattered Light, then the effective classification accuracy for this experiment exceeds 90%. It is also important to note that some misclassifications may arise from stochastic interactions between the

injected glitches and the specific noise realization into which they are injected, which can differ from the original conditions of the glitch sample. Additionally, labelling uncertainties in the Gravity Spy training set may contribute, and the discarded phase information in Q -scans may play a role—an issue we discuss further below.

2. Global morphological similarity via UMAP

To assess whether GlitchGAN captures the overall structure of the real glitch distribution, we project both real and synthetic glitch samples into a three-dimensional latent space using UMAP [40].

Figure 5 shows two complementary perspectives of this embedding space, computed from 500 samples per class. Corresponding individual class-wise UMAP projections in Figure 6 provide additional detail. Strong overlap between real (blue) and synthetic (green) samples indicates that GlitchGAN has learned the general morphological structure of the different glitch types.

For most classes (Blip, Fast Scattering, Scattered Light, Tomte), synthetic and real samples show close clustering with substantial overlap. The Whistle class exhibits the most noticeable divergence, with synthetic samples forming a partially separate cluster, suggesting that GlitchGAN captures the dominant Whistle morphology but does not fully explore the full morphological diversity within this class. This is consistent with the relative complexity of the Whistle waveform structure and may reflect limitations of the fixed GAN architecture when applied to this particular class without hyperparameter tuning.

3. Generating hybrid glitches via class interpolation

A key advantage of class-conditional generative models is the ability to produce “hybrid” or transitional morphologies by sampling mixed class vectors.

We explore two sampling strategies:

1. **Simplex mixing:** The class vector is normalized such that all coefficients sum to 1, i.e., sampled from a $k = 7$ simplex. This constrains mixing proportions to sum to 100%.
2. **Uniform mixing:** Each class coefficient is independently sampled from $[0, 1]$ without normalization constraint, allowing more extreme interpolations.

Figure 7 shows glitch examples from both mixing strategies, demonstrating intermediate morphologies. As shown, samples from these datasets exhibit mixed morphological characteristics drawn from multiple glitch classes as per the sampled class vector. The UMAP projection of mixed samples in Figure 8 confirms that these hybrid samples smoothly interpolate between pure-class

point clouds, with uniform mixing exploring the broadest region of the latent space.

While hybrid samples are not further analyzed in this study, they have potential applications for pipeline validation and data augmentation for glitch detection or classification systems.

C. Limitations of magnitude-only representations

We further evaluated the DiT model as a comparison; the full Gravity Spy and UMAP results are shown in Appendix C. DiT-generated samples overlap with real glitches primarily for low-frequency classes (e.g., Fast Scattering, Scattered Light), but show reduced diversity and separable clusters for most other classes where GlitchGAN was more successful.

A notable example is the Low Frequency Burst class: UMAP reveals two clearly distinct clusters (real vs. synthetic), yet Gravity Spy classifies 87% of DiT-generated samples correctly with 97% mean confidence. Figure 9 illustrates why: while the time-domain waveforms of real and DiT-generated Low Frequency Bursts are clearly distinct, their magnitude Q -scan representations appear similar, explaining the classifier’s high confidence despite the physically unrealistic time-domain morphology.

This finding underscores a critical limitation: classifiers operating on magnitude-only representations cannot reliably assess glitch realism. Complementary validation methods preserving phase information (such as UMAP on time-domain waveforms) are essential for ensuring synthetic glitches are suitable for detector simulations.

Overall, while the DiT model demonstrates that latent diffusion is a viable approach to conditional time-domain glitch generation, its ability to reproduce full morphological diversity—especially for mid- and high-frequency content—is limited relative to GlitchGAN.

V. DISCUSSION

A. Performance assessment and limitations

GlitchGAN successfully demonstrates time-domain synthesis of multiple realistic glitch classes. Substantial UMAP overlap between real and synthetic samples, combined with accurate Gravity Spy classification of generated glitches indicates morphological fidelity of generated glitches. However, there exists several limitations and caveats which we address here.

Partial failure was observed for the Whistle class, possibly due to insufficient regularization or architectural constraints on receptive field size. For example, smaller kernel sizes in generator and discriminator layers may be needed to capture the high-frequency content of Whistle glitches (please see Appendix A for details).




FIG. 4: Confusion matrix showing Gravity Spy classifications of glitches generated by GlitchGAN, with 100 samples per class. Each cell indicates the number of samples predicted for each class, with the average Gravity Spy confidence shown in brackets. The strong diagonal structure demonstrates that most synthetic glitches are correctly classified according to their intended class.




FIG. 5: Two complementary views of three-dimensional UMAP embeddings showing real and generated glitch samples. Each color corresponds to one of the seven glitch classes used during training.

Since GlitchGAN is trained on normalized waveforms, it does not learn an explicit SNR distribution, and thus generated glitches must be rescaled to a specified SNR for use in downstream applications. As shown in Section IV B 1, Gravity Spy classifications are strongly sen-

sitive to the injected SNR — consistent with the fact that individual glitch classes are partly characterized by their SNR distribution. Appropriate rescaling is therefore an important consideration for simulations and pipeline validation studies using these synthetic samples.




FIG. 6: Three-dimensional UMAP embeddings for each glitch class, comparing real (blue) and GlitchGAN-generated (green) samples. Each panel corresponds to one glitch type and is shown from a rotated viewpoint to highlight intra-class structure. The close overlap between real and synthetic distributions across most classes demonstrates that GlitchGAN effectively learns class-specific morphological features, though some variation remains for more complex glitch morphologies.




FIG. 7: Examples of hybrid glitch waveforms generated by GlitchGAN using two class-conditioning strategies. (a) Simplex mixing samples each class vector from a $k = 7$ simplex, ensuring coefficients sum to one. (b) Uniform mixing samples each coefficient independently from a uniform distribution in $U[0, 1]$. The corresponding sampled class vectors are shown above each plot, with the dominant glitch label also shown. Both methods produce hybrid morphologies interpolating between glitch classes, demonstrating GlitchGAN’s ability to explore smooth transitions across the learned latent class space.

Finally, we note that no hyperparameter optimization was performed for this specific dataset; the model is evaluated “out-of-the-box” from previous work (except for the dimensionality of the class vector), representing a conservative assessment of GlitchGAN’s potential. Tuning the parameters of the generator and discriminator architectures, as well as the training procedure, could further improve performance, particularly for the more challenging Whistle class.

B. Practical applications

The synthetic glitch population produced by GlitchGAN is immediately useful for:

1. **Mock data challenges (MDCs):** Enriching realistic detector simulations with synthetic instrumental noise
2. **Detector characterization:** Validating glitch detection, mitigation, and subtraction algorithms

3. **Training data augmentation:** Supplementing limited real glitch samples with GlitchGAN-generated samples

4. **Pipeline robustness testing:** Assessing pipeline sensitivity and false-alarm rates in the presence of instrumental noise morphologies

C. Recommendations for future work

There are several promising directions for future research building on this work. Firstly, hyperparameter optimization and architectural refinements could be performed to further improve the quality of generated data, particularly for the more challenging Whistle class. Architecture improvements may include the incorporation of attention mechanisms, smaller convolutional kernel sizes or dilated convolutions for multi-scale feature extraction.

Ultimately, the goal is to develop a comprehensive generative model capable of synthesizing the entire range of




FIG. 8: Three-dimensional UMAP embedding showing the distribution of class-mixed glitch samples relative to the vertex (pure-class) glitches. Vertex samples are coloured by their respective glitch class, while class-mixed samples generated using simplex and uniform class-conditioning vectors are shown in black. The smooth interpolation of mixed samples between class clusters illustrates how GlitchGAN captures continuous transitions across glitch morphologies in the latent class space, spanning the diversity of glitch types represented in the training set.

glitch morphologies in GW detectors with high fidelity, enabling robust simulations and pipeline validation for current and future GW detectors. Future work will iteratively include more glitch glitches, starting with more common and progressively incorporating less common ones.

Glitch characterization is a challenging task in GW data analysis. It would be beneficial to develop more sophisticated validation metrics that explicitly evaluate phase consistency, not just magnitude-level morphology, to ensure synthetic glitches are suitable for detector simulations. Such methods will be beneficial for evaluating the realism of synthetic glitches and guiding future improvements in generative modeling approaches.

While class conditional generation is a powerful tool, it may also be desirable to develop conditional generation strategies allowing control over glitch duration, SNR, and frequency content independently from class morphology. This would enable more flexible synthesis of glitches for specific applications, such as testing pipeline sensitivity to particular frequency bands or durations.

Finally, integrating these features with mock data challenge pipelines for rapid deployment in detector simulations will be an important step towards practical utility of these synthetic glitch populations for the GW data analysis community.

VI. CONCLUSIONS

We have demonstrated that a class-conditional deep generative model called GlitchGAN can effectively synthesize realistic GW detector glitches directly in the time

domain. By leveraging high-fidelity reconstructions from a glitch reconstruction framework called DeepExtractor and training a single unified GlitchGAN model based on the cDVGAN architecture across seven glitch classes, we achieve synthetic samples with strong morphological agreement with real data. GlitchGAN’s synthetic glitch population offers practical utility for mock data challenges, detector characterization studies, and validation of detection pipelines.

We evaluated both DeepExtractor’s reconstructions used to train GlitchGAN and GlitchGAN’s synthetic data using two complementary approaches: a supervised classification analysis using Gravity Spy, a state-of-the-art glitch classification algorithm trained on an extensive labeled dataset, and an unsupervised, data-driven UMAP analysis to visualize and compare the class distributions of real and synthetic glitches in a projected three-dimensional space.

The Gravity Spy classification results for DeepExtractor reconstructions showed that 89.57% of the reconstructed glitches were correctly classified, with a mean classifier confidence of 96.77% for correct predictions. Misclassified samples showed a lower confidence of 84.84%, suggesting that classifier confidence can serve as a quality veto. Overall, the misclassification patterns are generally explainable, typically misclassifying as morphologically similar classes, and do not indicate systematic failure of the reconstruction. The unsupervised UMAP analysis of DeepExtractor reconstructions revealed well-separated clusters according to the Gravity Spy labels, demonstrating that the reconstructions preserve class-specific morphological structure in the time domain.

For GlitchGAN’s generated glitches, Gravity Spy classified most samples correctly. Misclassifications primarily occurred for the Koi Fish and Scattered Light classes, which can largely be attributed to morphological similarities between similar classes and Gravity Spy’s classification being dependent on SNR. For instance, Scattered Light glitches were often misclassified as Fast Scattering, while Koi Fish are often misclassified as Blip or Tomte. These overlaps are expected, as the classes share similar time-frequency signatures. This is confirmed by increasing the injected SNR to 150, where 79% of generated Koi Fish samples are correctly classified.

The UMAP results showed that GlitchGAN-generated glitches generally occupy the same regions as real data in the embedded UMAP space, indicating that the model captures much of the morphological diversity within each glitch class. One exception is the Whistle class, which consists of the highest-frequency features in the dataset. For this class, synthetic samples formed a partially distinct cluster from the real data, suggesting that GlitchGAN did not fully capture the fine-scale, high-frequency structure characteristic of these glitches. Future work could address this by incorporating smaller convolutional kernels in the network (balanced against the need for long-range temporal modelling), dilated convolutions to better capture hierarchical structure, or attention mech-




FIG. 9: Two examples each of real (left) and DiT-generated (right) Low Frequency Burst glitches, showing both the time-domain waveform and the magnitude Q -scan representation. The time-domain waveforms are clearly distinct, yet the Q -scans appear similar—explaining why Gravity Spy classifies the DiT samples correctly with high confidence despite their physical unrealism.

anisms to enhance global feature representation.

Finally, this study highlights a fundamental limitation in using purely magnitude-based Q -scan representations for supervised glitch classification. We observed that a less robust diffusion-based model, which produces unrealistic Low Frequency Burst glitches in the time domain, can still yield correct Gravity Spy labels with high confidence. Although the time-domain waveforms of real and synthetic samples are clearly distinct, their Q -scan representations appear similar, explaining the classifier’s high confidence. This discrepancy is evident in the corresponding UMAP analysis of time-domain data, which reveals two distinct clusters separating real and synthetic samples. This finding highlights the importance of preserving phase information and integrating unsupervised analyses to complement glitch characterization efforts.

ACKNOWLEDGEMENTS

This research has made use of data or software obtained from the Gravitational Wave Open Science Center, a service of the LIGO Scientific Collaboration, the Virgo Collaboration, and KAGRA. This material is based upon work supported by NSF’s LIGO Laboratory which is a major facility fully funded by the National Science Foundation, as well as the Science and Technology Facilities Council (STFC) of the United Kingdom, the Max-Planck-Society (MPS), and the State of Niedersachsen/Germany for support of the construction of Advanced LIGO and construction and operation of the GEO600 detector. Additional support for Advanced LIGO was provided by the Australian Research Council. Virgo is funded, through the European Gravitational Observatory (EGO), by the French Centre National de Recherche Scientifique (CNRS), the Italian Istituto Nazionale di Fisica Nucleare (INFN) and the Dutch Nikhef, with contributions by institutions from Belgium, Germany, Greece, Hungary, Ireland, Japan, Monaco, Poland, Portugal, Spain. KAGRA is supported by Ministry of Education, Culture, Sports, Science and Technology (MEXT), Japan Society for the Promotion of Science (JSPS) in Japan; National Research Foundation (NRF) and Ministry of Science and ICT (MSIT) in Korea; Academia Sinica (AS) and National Science and Technology Council (NSTC) in Taiwan. The authors are grateful for computational resources provided by the LIGO Laboratory and supported by the National

Science Foundation Grants No. PHY-0757058 and No. PHY-0823459.

Appendix A: cDVGAN Architecture

Table IV shows the cDVGAN architecture featured in GlitchGAN.

		Discriminator (3.5M parameters)			
Operation	Output shape	Kernel size	Stride	Dropout	Activation
Input	(1024)	—	—	0	—
Reshape	(64,16)	—	—	0	—
Convolutional	(64,128)	14	2	0.5	Leaky ReLU
Convolutional	(32,128)	14	2	0.5	Leaky ReLU
Convolutional	(16,256)	14	2	0.5	Leaky ReLU
Convolutional	(8,256)	14	2	0.5	Leaky ReLU
Convolutional	(4,512)	14	2	0.5	Leaky ReLU
Global Avg. Pooling	(512)	—	—	0.5	—
Avg. Pooling Dense	(128)	—	—	0.2	Leaky ReLU
Dense	(1)	—	—	0	Linear
Class input	(7)	—	—	—	—
Class Dense	(128)	—	—	0	Linear
Scalar product	(1)	—	—	—	—
Dense + Scalar product	(1)	—	—	—	—

		DV Discriminator (1.1M parameters)			
Operation	Output shape	Kernel size	Stride	Dropout	Activation
Input	(1023)	—	—	0	—
Dense	(512)	—	—	0	Leaky ReLU
Reshape	(32,16)	—	—	0	—
Convolutional	(32,64)	5	2	0.5	Leaky ReLU
Convolutional	(16,128)	5	2	0.5	Leaky ReLU
Convolutional	(8,256)	5	2	0.5	Leaky ReLU
Convolutional	(4,256)	5	2	0.5	Leaky ReLU
Global Avg. Pooling	(256)	—	—	0.5	—
Avg. Pooling Dense	(128)	—	—	0.2	Leaky ReLU
Dense	(1)	—	—	0	Linear
Class input	(7)	—	—	—	—
Class Dense	(128)	—	—	0	Linear
Scalar product	(1)	—	—	—	—
Dense + Scalar product	(1)	—	—	—	—

		Generator (3.5M parameters)			
Operation	Output shape	Kernel size	Stride	BN	Activation
Latent input	(100)	—	—	×	—
Class input	(7)	—	—	×	—
Class Dense	(32)	—	—	×	—
Concatenate	(132)	—	—	×	—
Dense	(1024)	—	—	×	ReLU
Reshape	(32,32)	—	—	×	—
Transposed conv.	(64,512)	18	2	✓	ReLU
Transposed conv.	(128,256)	18	2	✓	ReLU
Transposed conv.	(256,128)	18	2	✓	ReLU
Transposed conv.	(512,64)	18	2	✓	ReLU
Transposed conv.	(1024,1)	18	2	×	Linear
Flatten	(1024)	—	—	×	—
Optimizer	RMSprop ($\alpha = 0.0001$)				
Batch size	512				
Epochs	500				
Loss	Wasserstein				

TABLE IV: Architecture and hyperparameters of cDVGAN featured in GlitchGAN. The model includes a base discriminator, derivative (DV) discriminator, and generator.

Appendix B: cDVGAN Loss

Figure 10 shows the training loss for each individual model component.

Appendix C: DiT evaluation results

Figures 11 and 12 show the full UMAP and Gravity Spy evaluation results for the DiT comparison model discussed in Section IV.




FIG. 10: *GlitchGAN* training loss over 500 epochs.



FIG. 11: *Gravity Spy* confusion matrix for samples generated by the DiT model (100 per class). Each cell shows the number of predictions and mean classifier confidence in brackets. Despite the poor UMAP overlap, *Gravity Spy* achieves high confidence classification for most classes, illustrating the limitation of magnitude-only validation.

- [1] LIGO Scientific Collaboration and others, Advanced LIGO, *Class. Quant. Grav.* **32**, 074001 (2015), [arXiv:1411.4547 \[gr-qc\]](#).
- [2] F. Acernese *et al.* (VIRGO), Advanced Virgo: a second-generation interferometric gravitational wave detector, *Class. Quant. Grav.* **32**, 024001 (2015), [arXiv:1408.3978 \[gr-qc\]](#).
- [3] R. Abbott *et al.*, Gwtc-3: Compact binary coalescences observed by ligo and virgo during the second part of the third observing run, *Physical Review X* **13**, 10.1103/physrevx.13.041039 (2023).
- [4] The LIGO Scientific Collaboration *et al.*, Characterization of the ligo detectors during their third observing run, *Classical and Quantum Gravity* **38**, 135003 (2021).
- [5] D. Davis *et al.*, Ligo detector characterization in the second and third observing runs, *Classical and Quantum Gravity* **38**, 135014 (2021).
- [6] F. Acernese *et al.* (Virgo), Virgo detector characterization and data quality: results from the O3 run, *Class. Quant. Grav.* **40**, 185006 (2023), [arXiv:2210.15633 \[gr-qc\]](#).
- [7] T. Akutsu *et al.*, Overview of KAGRA: Calibration, detector characterization, physical environmental monitors, and the geophysics interferometer, *Progress of Theoretical and Experimental Physics* **2021**, 05A102 (2021), <https://academic.oup.com/ptep/article-pdf/2021/5/05A102/38109702/ptab018.pdf>.
- [8] B. P. Abbott *et al.* (LIGO Scientific, Virgo), Effects of data quality vetoes on a search for compact binary coalescences in Advanced LIGO's first observing run, *Class. Quant. Grav.* **35**, 065010 (2018), [arXiv:1710.02185 \[gr-qc\]](#).
- [9] L. Blackburn *et al.*, The LSC Glitch Group: Monitoring Noise Transients during the fifth LIGO Science Run, *Class. Quant. Grav.* **25**, 184004 (2008), [arXiv:0804.0800 \[gr-qc\]](#).
- [10] B. P. Abbott *et al.* (LIGO Scientific, Virgo), Characterization of transient noise in Advanced LIGO relevant to gravitational wave signal GW150914, *Class. Quant. Grav.* **33**, 134001 (2016), [arXiv:1602.03844 \[gr-qc\]](#).
- [11] C. Pankow *et al.*, Mitigation of the instrumental noise transient in gravitational-wave data surrounding gw170817, *Phys. Rev. D* **98**, 084016 (2018).




FIG. 12: Three-dimensional UMAP projection of real (orange) and DiT-generated (blue) glitches across all seven classes. The large panel shows the overall structure, while smaller panels display each class individually. Overlapping clusters indicate morphological similarity between real and synthetic samples; separable clusters for most classes indicate that DiT does not reproduce the full morphological structure of real glitches, in contrast to GlitchGAN (Figure 5).

- [12] E. Payne, S. Hourihane, J. Golomb, R. Udall, D. Davis, and K. Chatziioannou, Curious case of gw200129: Interplay between spin-precession inference and data-quality issues, *Phys. Rev. D* **106**, 104017 (2022).
- [13] A. Gupta *et al.*, Possible causes of false general relativity violations in gravitational wave observations (2024), [arXiv:2405.02197 \[gr-qc\]](https://arxiv.org/abs/2405.02197).
- [14] B. P. Abbott *et al.* (LIGO Scientific Collaboration and Virgo Collaboration), Gw170817: Observation of gravitational waves from a binary neutron star inspiral, *Phys. Rev. Lett.* **119**, 161101 (2017).
- [15] B. P. Abbott *et al.*, Gw190425: Observation of a compact binary coalescence with total mass $\sim 3.4 m_{\odot}$, *The Astrophysical Journal Letters* **892**, L3 (2020).
- [16] R. Abbott *et al.*, Gw190814: Gravitational waves from the coalescence of a 23 solar mass black hole with a 2.6 solar mass compact object, *The Astrophysical Journal Letters* **896**, L44 (2020).
- [17] R. Abbott *et al.*, Observation of gravitational waves from two neutron star–black hole coalescences, *The Astrophysical Journal Letters* **915**, L5 (2021).
- [18] S. Hild, M. Abernathy, F. Acernese, P. Amaro-Seoane, N. Andersson, K. Arun, F. Barone, B. Barr, M. Barsuglia, and M. B. e, Sensitivity studies for third-generation gravitational wave observatories, *Class. Quant. Grav.* **28**, 094013 (2011).
- [19] E. D. Hall, Cosmic explorer: A next-generation ground-based gravitational-wave observatory, *Galaxies* **10**, 10.3390/galaxies10040090 (2022).
- [20] B. P. Abbott *et al.*, Prospects for observing and localizing gravitational-wave transients with advanced ligo, advanced virgo and kagra, *Living Reviews in Relativity* **23**, 10.1007/s41114-020-00026-9 (2020).
- [21] J. Abadie *et al.*, Search for gravitational waves from compact binary coalescence in ligo and virgo data from s5 and vsr1, *Physical Review D* **82**, 10.1103/physrevd.82.102001 (2010).
- [22] J. Abadie *et al.*, All-sky search for gravitational-wave bursts in the first joint ligo-geo-virgo run, *Physical Review D* **81**, 10.1103/physrevd.81.102001 (2010).
- [23] C. Biwer *et al.*, Validating gravitational-wave detections: The advanced ligo hardware injection system, *Phys. Rev. D* **95**, 10.1103/physrevd.95.062002 (2017).
- [24] A. E. Tolley *et al.*, ArchEnemy: removing scattered-light glitches from gravitational wave data, *Class. Quant. Grav.* **40**, 165005 (2023), [arXiv:2301.10491 \[gr-qc\]](https://arxiv.org/abs/2301.10491).
- [25] R. Bondarescu, A. Lundgren, and R. Macas, Quasi-physical model for removing short glitches from LIGO

- and Virgo data, *Phys. Rev. D* **108**, 122004 (2023), [arXiv:2309.06594 \[gr-qc\]](#).
- [26] C.-H. Liao and F.-L. Lin, Deep generative models of gravitational waveforms via conditional autoencoder, *Phys. Rev. D* **103**, 124051 (2021), [arXiv:2101.06685 \[astro-ph.IM\]](#).
- [27] T. Eccleston and M. C. Edwards, Generative adversarial network for stellar core-collapse gravitational waves, *Phys. Rev. D* **110**, 104055 (2024).
- [28] I. J. Goodfellow *et al.*, Generative Adversarial Networks (2014), [arXiv:1406.2661 \[stat.ML\]](#).
- [29] J. Yan, A. P. Leung, and D. C. Y. Hui, [On improving the performance of glitch classification for gravitational wave detection by using generative adversarial networks](#) (2022).
- [30] J. Powell, L. Sun, K. Gereb, P. D. Lasky, and M. Dollmann, Generating transient noise artefacts in gravitational-wave detector data with generative adversarial networks, *Classical and Quantum Gravity* **40**, 035006 (2023).
- [31] M. Lopez *et al.*, [Simulating transient noise bursts in ligo with generative adversarial networks](#) (2022).
- [32] M. Lopez, V. Boudart, S. Schmidt, and S. Caudill, [Simulating transient noise bursts in ligo with gengli](#) (2022).
- [33] N. J. Cornish and T. B. Littenberg, Bayeswave: Bayesian inference for gravitational wave bursts and instrument glitches, *Class. Quant. Grav.* **32**, 135012 (2015).
- [34] J. McGinn *et al.*, Generalised gravitational wave burst generation with generative adversarial networks, *Class. Quant. Grav.* **38**, 155005 (2021).
- [35] T. Dooney *et al.*, Time-domain reconstruction of signals and glitches in gravitational wave data with deep learning, *Physical Review D* **10.1103/s91m-c2jw** (2025).
- [36] T. Dooney, R. L. Curier, D. S. Tan, M. Lopez, C. Van Den Broeck, and S. Bromuri, One flexible model for multiclass gravitational wave signal and glitch generation, *Phys. Rev. D* **110**, 022004 (2024).
- [37] T. Dooney, S. Bromuri, and L. Curier, Dvgan: Stabilize Wasserstein GAN training for time-domain gravitational wave physics, in *2022 IEEE International Conference on Big Data (Big Data)* (Osaka, Japan, 2022) pp. 5468–5477.
- [38] M. Zevin *et al.*, Gravity spy: integrating advanced LIGO detector characterization, machine learning, and citizen science, *Class. Quant. Grav.* **34**, 064003 (2017).
- [39] Y. Wu, M. Zevin, C. P. L. Berry, K. Crowston, C. Østerlund, Z. Doctor, S. Banagiri, C. B. Jackson, V. Kalogera, and A. K. Katsaggelos, Advancing glitch classification in gravity spy: Multi-view fusion with attention-based machine learning for advanced ligo’s fourth observing run (2024), [arXiv:2401.12913 \[gr-qc\]](#).
- [40] L. McInnes, J. Healy, and J. Melville, UMAP: Uniform Manifold Approximation and Projection for Dimension Reduction, *arXiv preprint* (2018), [arXiv:1802.03426](#).
- [41] J. Sohl-Dickstein, E. A. Weiss, N. Maheswaranathan, and S. Ganguli, [Deep unsupervised learning using nonequilibrium thermodynamics](#) (2015).
- [42] M. Cabero *et al.*, Blip glitches in Advanced LIGO data, *Class. Quant. Grav.* **36**, 15 (2019), [arXiv:1901.05093 \[physics.ins-det\]](#).
- [43] S. Soni *et al.*, Discovering features in gravitational-wave data through detector characterization, citizen science and machine learning, *Class. Quant. Grav.* **38**, 195016 (2021), [arXiv:2103.12104 \[gr-qc\]](#).
- [44] L. Nuttall *et al.*, Improving the Data Quality of Advanced LIGO Based on Early Engineering Run Results, *Class. Quant. Grav.* **32**, 245005 (2015), [arXiv:1508.07316 \[gr-qc\]](#).
- [45] M. Arjovsky, S. Chintala, and L. Bottou, Wasserstein generative adversarial networks, in *Proceedings of the 34th International Conference on Machine Learning*, Proceedings of Machine Learning Research, Vol. 70, edited by D. Precup and Y. W. Teh (PMLR, 2017) pp. 214–223.
- [46] I. Gulrajani, F. Ahmed, M. Arjovsky, V. Dumoulin, and A. Courville, [Improved training of wasserstein gans](#) (2017).
- [47] M. Mirza and S. Osindero, Conditional generative adversarial nets, *arXiv preprint* (2014), [arXiv:1411.1784](#).
- [48] M. Abadi *et al.*, [TensorFlow: Large-scale machine learning on heterogeneous systems](#) (2015), software available from tensorflow.org.
- [49] J. Ho, A. Jain, and P. Abbeel, [Denoising diffusion probabilistic models](#) (2020), [arXiv:2006.11239 \[cs.LG\]](#).
- [50] J. Ho and T. Salimans, Classifier-free diffusion guidance (2022), [arXiv:2207.12598 \[cs.LG\]](#).
- [51] T. T. Cai and R. Ma, [Theoretical foundations of t-sne for visualizing high-dimensional clustered data](#) (2022), [arXiv:2105.07536 \[stat.ML\]](#).
- [52] L. McInnes, J. Healy, and J. Melville, [Umap: Uniform manifold approximation and projection for dimension reduction](#) (2020), [arXiv:1802.03426 \[stat.ML\]](#).
- [53] J. Areeda, Are blips and koi fish related?, <https://blog.gravityspy.org/2017/10/19/are-blips-and-koi-fish-related/> (2017), accessed on 11-06-2025.